

Yincheng Zhou

(267) 788-2488 | yzhou29@seas.upenn.edu | github.com/ArtysicistZ | artysicistz.github.io

EDUCATION

University of Pennsylvania

Vagelos Integrated Program in Energy Research (VIPER)

- **Dual Degree:** B.S.E. in Computer Science (AI Concentration) | B.A. in Physics & Astronomy
- **Relevant Coursework:** Big Data Analysis, Computer Systems, Data Structures & Algorithms, OOP
- **GPA:** 4.0 / 4.0

Philadelphia, PA
September 2025 – Present

EXPERIENCE

GRASP Laboratory | University of Pennsylvania [\[GitHub\]](#)

Machine Learning Researcher, Advisor: Prof. Pratik Chaudhari

- Proposed SPD, a sparse monocular depth estimator that predicts depth only at queried pixels, achieving $6\times$ inference speedup over dense methods while preserving accuracy (AbsRel 0.11 on NYU Depth V2) for real-time robotics.
- Developed a dense-train / sparse-infer paradigm with a 350K-param split decoder and cross-attention fusion on DINOv2 features, proving sparse output matches dense pixel-wise prediction through systematic ablation across 15 architecture variants.

Philadelphia, PA

January 2026 – Present

Welta, Inc. | AI-powered cross-platform ad optimization engine [\[Demo\]](#)

Software Engineer Intern

- Designed ad spend optimizer using Bayesian response curves and convex optimization to reallocate budget across channels at fixed cost; validated on 93 e-commerce brands, delivering 23% higher return on ad spend versus manual allocation.
- Engineered ad simulation engine that tokenizes creatives and models persona-level reactions via LLM; reduced inference latency to 20% of baseline while improving prediction accuracy by 55%.

Remote

March 2026 – Present

Human Machine Intelligence Laboratory | Peking University [\[HuggingFace\]](#)

Machine Learning Researcher, Advisor: Prof. Shanghang Zhang

- Developed MCTS-guided SFT for GUI agent training that collects and trains on step-level decisions with difficulty weighting, outperforming trajectory-level RL (GRPO) at 25% compute; raised UI-TARS-7B's OSWorld success from 12.3% to 17.7%.
- Engineered distributed rollout infrastructure across 112 Docker VMs and 8 A100 GPUs with fault-tolerant container monitoring and action-replay VM cloning; collected 2,927 trajectories and curated 35K steps via 200 parallel LLM auditing agents.

Remote

March 2026 – Present

Franklink, Inc. | World's first AI-native networking platform [\[Demo\]](#)

Founding Software Engineer

- Designed iMessage-based multi-agent orchestration layer with persisted memory, selective dispatch across 5 sub-agents, and agentic RAG over 8 permissioned tools (DB, email, calendar); reduced user-facing latency by 30%.
- Engineered a Kafka-backed event pipeline with at-least-once delivery, idempotency keys, and retry + DLQ; sustained 50+ concurrent sessions with p99 latency within 40% of single-session baseline; deployed on AWS ECS.

Philadelphia, PA

October 2025 – February 2026

PROJECTS

AlphaOne | Subject-oriented financial sentiment platform [\[GitHub\]](#) [\[Demo\]](#)

- Fine-tuned DeBERTa for per-ticker aspect-based sentiment in multi-entity financial text, achieving 82.5% accuracy (+65pp over FinBERT); curated 6,287 training triples via Reddit scraping, LLM labeling, and manual audit of 4,501 pairs.
- Engineered a 3-pass batch NLP pipeline (split, infer, commit) with claim-based row locking for concurrent-safe ingestion across 7 subreddits; serves real-time per-ticker trends through a full-stack dashboard.

August 2025 – February 2026

F-Zero Racing Agent | RL agent for competitive SNES racing [\[GitHub\]](#)

- Trained a PPO agent across 80 parallel SNES emulators with custom reward shaping to race F-Zero, reaching global top 300 on a task previously deemed unsolvable due to sparse terminal-only rewards over 3,000+ step horizons.
- Engineered dense reward from raw RAM by spline-interpolating 58 checkpoints into a 660-point track centerline; designed dual-input CNN+MLP policy fusing 4-frame visual input with 59-dim state features.

February 2026 – April 2026

IncidentGuard | Multi-agent incident remediation system [\[GitHub\]](#)

- Built a multi-agent system that predicts remediation outcomes before execution, rejecting dangerous reflexes via counterfactual ranking; reduced simulated recovery latency by 46% over naive baseline across 4 failure scenarios.
- Implemented closed-loop plan-execute-verify pipeline with live Prometheus metrics, OpenClaw browser automation, and DuckDB case persistence; system gracefully degrades across 3 fallback paths when infrastructure is unavailable.

March 2026 – April 2026

NL2SQL Bot | Agentic NL-to-SQL orchestration [\[GitHub\]](#)

- Orchestrated 5 specialized agents on Google ADK with loop-guarded retry, per-agent model routing, and automated chart generation across 6 visualization types from natural language queries.
- Enforced read-only execution via 18 blocked SQL patterns and table allowlisting; abstracted dialect-specific rules across MySQL, PostgreSQL, and SQLite for cross-database portability.

December 2025 – January 2026

SKILLS & ACHIEVEMENTS

Languages: C++, Python, Java, JavaScript/TypeScript, SQL, Assembly, OCaml

Backend & Infrastructure: FastAPI, Spring Boot, Kafka, Redis, Celery, Docker, AWS, Ray, Linux, Git

ML & Data: PyTorch, CUDA, vLLM, HuggingFace Transformers, spaCy, Pandas, LangGraph

Achievements: Team China, International Science and Engineering Fair (ISEF); Silver Medal, Team China, International Young Physicists' Tournament (IYPT); Second Prize, China Mathematics Olympiad; Second Prize, China Physics Olympiad